

La práctica de los experimentos de encuesta: diseños y aplicaciones recientes

Santiago López-Cariboni, Universidad de la República - dECON

Activity!

1. Ask you to guess a number
2. Number off 1 and 2 across the room
3. Group 2, close your eyes

Activity!

Group 1

Think about whether the population of Chicago is more or less than 500,000 people.

What do you think the population of Chicago is?

Activity!

1. Ask you to guess a number
2. Number off 1 and 2 across the room
3. Group 2, close your eyes
4. Group 1, close your eyes

Activity!

Group 2

Think about whether the population of Chicago is more or less than 10,000,000 people.

What do you think the population of Chicago is?

Enter your data

Go here: <https://bit.ly/2Sxsp17>

Enter your guess and your group number

Results

- ▶ True population: 2.79 million
- ▶ What did you guess? (See Responses)
- ▶ What's going on here?
 - ▶ An experiment!
 - ▶ Demonstrates “anchoring” heuristic
- ▶ Experiments are easy to analyze, but only if designed and implemented well

1. Introductions
2. Course Outline
3. History and Logic

Santiago López Cariboni

Profesor asociado, Departamento de Economía, UdelaR

2008–14: PhD Univ. of Essex (UK)

Intereses: CPE, IPE

Email: santiago.lopez@cienciassociales.edu.uy

Who are you?

- ▶ Introduce yourself
- ▶ Where are you from?
- ▶ What do you hope to learn from the course?

Quick Survey

- ▶ How many of you have worked with survey data before?
- ▶ Of those, how many of you have performed a survey before?
- ▶ How many of you have worked with experimental data before?
- ▶ Of those, how many of you have performed an experiment before?

Course Outline

Course Materials

Course is available at:

<https://eva.eduper.udelar.edu.uy/>

Learning Outcomes

By the end of the course, you should be able to:

1. Explain how to analyze experiments quantitatively.
2. Explain how to design experiments that speak to relevant research questions and theories.
3. Evaluate the uses and limitations of several common survey experimental paradigms.
4. Identify practical issues that arise in the implementation of experiments and evaluate how to anticipate and respond to them.

Programa de sesiones por tema

1. Clase 01/11: Historia y contexto reciente de los experimentos de encuesta
2. Clase 03/11 - Clase 05/11 - : Tipos de experimentos de encuesta
3. Clase 08/11: La implementación de experimentos de encuesta
4. Clase 10/11 - Clase 12/11 - : Análisis estadístico de experimentos
5. Clase 15/11: Revisión y Presentación de proyectos

History and Logic

Experiments: History I

Oxford English Dictionary defines “experiment” as:

- ▶ A scientific procedure undertaken to make a discovery, test a hypothesis, or demonstrate a known fact
- ▶ A course of action tentatively adopted without being sure of the outcome

Experiments: History II

-“Experiments” have a very long history

- ▶ Major advances in design and analysis of experiments based on agricultural and later biostatistical research in the 19th century (Fisher, Neyman, Pearson, etc.)
- ▶ Multiple origins in the social sciences
 - ▶ First randomized experiment by Peirce and Jastrow (1884)
 - ▶ Gosnell (1924)
 - ▶ LaLonde (1986)
 - ▶ Gerber and Green (2000)

Experiments: History III

- ▶ Rise of surveys in the behavioral revolution
 - ▶ Survey research not heavily experimental because interviewing was mostly paper-based
 - ▶ Split ballots (e.g., Schuman & Presser; Bishop)
- ▶ 1983: Merrill Shanks and the Berkeley Survey Research Center develop CATI
- ▶ Mid-1980s: Paul Sniderman & Tom Piazza performed the first modern survey experiment¹
 - ▶ Then: the “first multi-investigator”
 - ▶ Later: Arthur Lupia and Diana Mutz created TESS

¹Sniderman, Paul M., and Thomas Piazza. 1993. *The Scar of Race*. Cambridge, MA: Harvard University Press.

TESS

- ▶ Time-Sharing Experiments for the Social Sciences
- ▶ Multi-disciplinary initiative that provides infrastructure for survey experiments on nationally representative samples of the United States population
- ▶ Great resource for survey experimental materials, designs, and data
- ▶ Funded by the U.S. National Science Foundation
- ▶ Anyone anywhere in the world can apply
- ▶ See also: LISS, Bergen's Citizen Panel, Gothenburg's Citizen Panel

[Home](#)
[FAQ](#)
[Introduction](#)
[How to Submit](#)
[Short Studies](#)
[Study Archives & Data](#)
[Past Announcements](#)
[Submit a Proposal](#)
[Contact Us](#)

[2020](#) | [2019](#) | [2018](#) | [2017](#) | [2016](#) | [2015](#) | [2014](#) | [2013](#) | [2012](#) | [2011](#) | [2010](#) | [2009](#) | [2008](#) | [2007](#) | [2006](#) | [2005](#) | [2004](#) | [2003](#)

Past TESS Studies

Note: Because authors of successful TESS proposals have exclusive right to TESS data for one year after data are delivered to them, only studies fielded more than one year ago are listed here.

2020

Polarization in Pandemics: Partisan Differences in Coronavirus Social Distancing

Hunt Allcott, Microsoft Research

Mechanical Asians and Animalistic Blacks: The Political Implications of The Symmetry of Two Forms of Dehumanization in Racial Perceptions

Hui Bai, University of Minnesota - Twin Cities

Lotteries: A Behavioral Welfare Evaluation

Hunt Allcott, Microsoft Research
Benjamin Lockwood, University of Pennsylvania
Dmitry Taubinsky, University of California, Berkeley

How Candidates Can Signal Sincerity in an Era of Cynicism

Scott Clifford, University of Houston
Elizabeth Simas, University of Houston

Everyone's Doing It: Affective Polarization is Inflated by Social Pressure

Elizabeth C. Connors, University of South Carolina

A Conditional Commitment? Partisan Identity and Support for Democracy in the United States

Matthew H. Graham, Yale University



The First Survey Experiment?

Hadley Cantril (1940) asks 3000 Americans either:
Do you think the U.S. should do
more than it is now doing to
help England and France?

- ▶ Yes
- ▶ No

The First Survey Experiment?

Hadley Cantril (1940) asks 3000 Americans either:

Do you think the U.S. should do more than it is now doing to help England and France?

- ▶ Yes
- ▶ No

Do you think the U.S. should do more than it is now doing to help England and France **in their fight against Hitler?**

- ▶ Yes
- ▶ No

The First Survey Experiment?

Hadley Cantril (1940) asks 3000 Americans either:

Do you think the U.S. should do more than it is now doing to help England and France?

- ▶ Yes 13%
- ▶ No

Do you think the U.S. should do more than it is now doing to help England and France **in their fight against Hitler?**

- ▶ Yes 22%
- ▶ No

The “Hitler effect” was $22\% - 13\% = 9\%$

Definitions I

- ▶ A randomized experiment is:

The observation of units after, and possibly before, a randomly assigned intervention in a controlled setting, which tests one or more precise causal expectations

- ▶ If we manipulate the thing we want to know the effect of (X), and control (i.e., hold constant) everything we do not want to know the effect of (\mathbf{Z}), the only thing that can affect the outcome (Y) is X .

Definitions II

- ▶ A survey experiment is just an experiment that occurs in a survey context
 - ▶ As opposed to in the field or in a laboratory
- ▶ Can be in any mode (face-to-face, CATI, IVR, CASI, etc.)
- ▶ May or may not involve a representative population Mutz (2011): “population-based survey experiments”

Definitions II

Unit: A physical object at a particular point in time

Definitions II

Treatment: An intervention, whose effect(s) we wish to assess relative to some other (non-)intervention

Synonyms: manipulation, intervention, factor, condition, cell

Definitions II

Outcome: The variable we are trying to explain

Definitions II

Potential outcomes: The outcome value for each unit that we would observe if that unit received each treatment

Multiple potential outcomes for each unit, but we only observe one of them

Definitions II

Causal effect: The comparisons between the unit-level potential outcomes under each intervention

This is what we want to know!

Definitions II

Average causal effect: Difference in mean outcomes between treatment groups

This is almost what we want to know!

Example

Unit: Americans in 1940

Outcome: Support for military intervention

Treatment: Mentioning Hitler versus not

Potential outcomes :

1. Support in “Hitler” condition
2. Support in control condition

Causal effect: Difference in support between the two question wordings for each respondent

- ▶ Individual treatment effect not observable!
- ▶ Average effect (ATE) is the mean-difference

Questions?

Why are experiments useful?

Causal inference!

Addressing Confounding

In observational research. . .

1. Correlate a “putative” cause (X) and an outcome (Y), where X temporally precedes Y

Addressing Confounding

In observational research. . .

1. Correlate a “putative” cause (X) and an outcome (Y), where X temporally precedes Y
2. Identify all possible confounds (Z)

Addressing Confounding

In observational research. . .

1. Correlate a “putative” cause (X) and an outcome (Y), where X temporally precedes Y
2. Identify all possible confounds (Z)
3. “Condition” on all confounds

Addressing Confounding

In observational research. . .

1. Correlate a “putative” cause (X) and an outcome (Y), where X temporally precedes Y
2. Identify all possible confounds (Z)
3. “Condition” on all confounds
 - ▶ Calculate correlation between X and Y at each combination of levels of Z

Addressing Confounding

In observational research. . .

1. Correlate a “putative” cause (X) and an outcome (Y), where X temporally precedes Y
2. Identify all possible confounds (\mathbf{Z})
3. “Condition” on all confounds
 - ▶ Calculate correlation between X and Y at each combination of levels of \mathbf{Z}
4. Basically: $Y = \beta_0 + \beta_1 X + \beta_{2-k} \mathbf{Z} + \epsilon$

Media
Coverage



Demographics

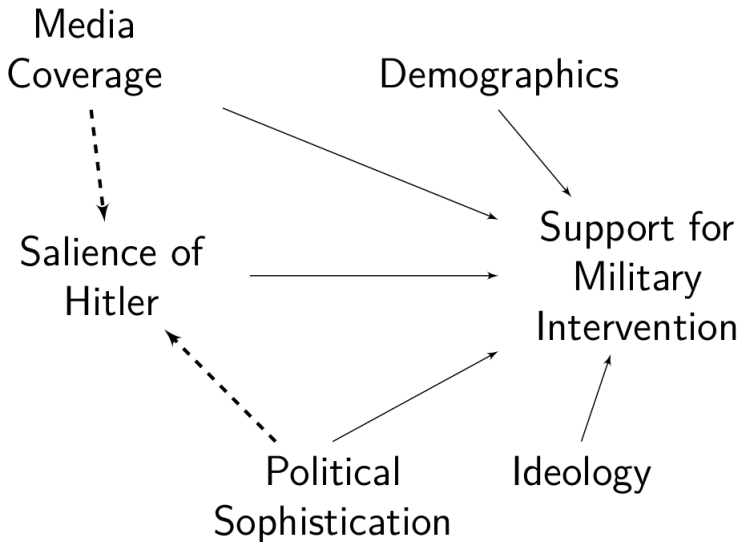
Salience of
Hitler

Support for
Military
Intervention

Political
Sophistication

Ideology





Experiments are different

1. Causal inferences from design not analysis
2. Solves both temporal ordering and confounding
 - ▶ Treatment (X) applied by researcher before outcome (Y)
 - ▶ Randomization eliminates confounding (Z)
 - ▶ We don't need to "control" for anything
3. Basically: $Y = \beta_0 + \beta_1 X + \epsilon$
4. Thus experiments are a "gold standard"

Mill's Method of Difference

If an instance in which the phenomenon under investigation occurs, and an instance in which it does not occur, have every circumstance save one in common, that one occurring only in the former; the circumstance in which alone the two instances differ, is the effect, or cause, or an necessary part of the cause, of the phenomenon.

Mill's Method of Difference

If an instance in which the phenomenon under investigation occurs, and an instance in which it does not occur, **have every circumstance save one in common**, that one occurring only in the former; **the circumstance in which alone the two instances differ, is the effect, or cause, or an necessary part of the cause, of the phenomenon.**

Questions?

Neyman-Rubin Potential Outcomes Framework

If we are interested in some outcome Y , then for every unit i , there are numerous “potential outcomes” Y^* only one of which is visible in a given reality. Comparisons of (partially unobservable) potential outcomes indicate causality.

Neyman-Rubin Potential Outcomes Framework

Concisely, we typically discuss two potential outcomes:

- ▶ Y_{0i} , the potential outcome realized if $X_i = 0$
 - ▶ (because $D_i = 0$, assigned to control)
- ▶ Y_{1i} , the potential outcome realized if $X_i = 1$
 - ▶ (because $D_i = 1$, assigned to treatment)

Experimental Inference I

Each unit has multiple *potential* outcomes, but we only observe one of them, randomly

Experimental Inference I

Each unit has multiple *potential* outcomes, but we only observe one of them, randomly

In this sense, we are sampling potential outcomes from each unit's population of potential outcomes

unit	low	high	control	etc.
1	?	?	?	...
2	?	?	?	...
3	?	?	?	...
4	?	?	?	...

Experimental Inference II

- ▶ We cannot see individual-level causal effects
- ▶ We can see average causal effects
 - ▶ Ex.: Average difference in military support among those thinking of Hitler versus not
- ▶ We want to know: $TE_i = Y_{1i} - Y_{0i}$

Experimental Inference III

- ▶ We want to know: $TE_i = Y_{1i} - Y_{0i}$ for every i in the population
- ▶ We can average: $E[TE_i] = E[Y_{1i} - Y_{0i}] = E[Y_{1i}] - E[Y_{0i}]$
- ▶ But we still only see one potential outcome for each unit:
 $ATE_{naive} = E[Y_1|X = 1] - E[Y_0|X = 0]$
- ▶ *Is this what we want to know?*

Experimental Inference IV

- ▶ What we want and what we have:

$$ATE = E[Y_1] - E[Y_0]$$

$$ATE_{naive} = E[Y_1|X = 1] - E[Y_0|X = 0]$$

Are the following statements true?

$$E[Y_1] = E[Y_1|X = 1]$$

$$E[Y_0] = E[Y_0|X = 0]$$

- ▶ Not in general

Experimental Inference V

- ▶ Only true when both of the following hold:

$$E[Y_1] = E[Y_1|X = 1] = E[Y_1|X = 0]$$

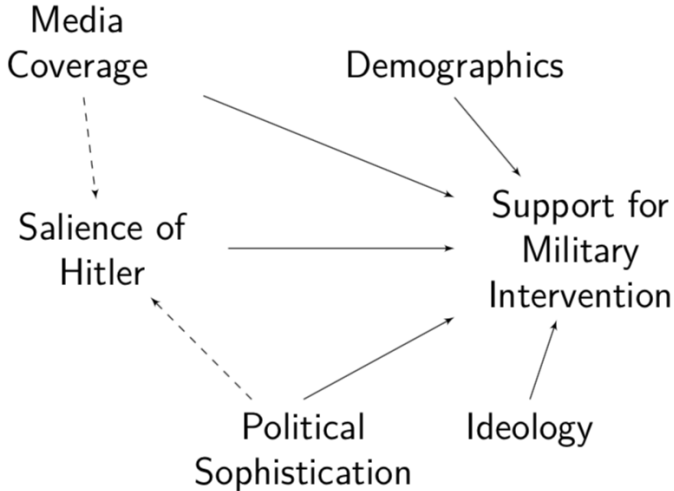
$$E[Y_0] = E[Y_0|X = 1] = E[Y_0|X = 0]$$

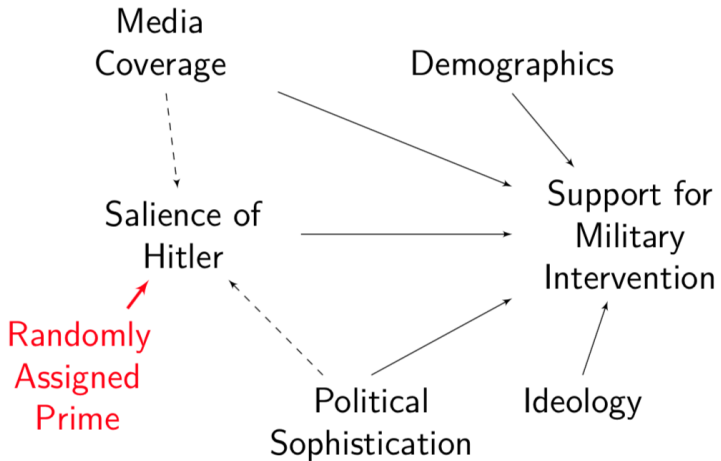
- ▶ In that case, potential outcomes are independent of treatment assignment
- ▶ If true (e.g., due to randomization of X), then:

$$\begin{aligned}ATE_{naive} &= E[Y_1|X = 1] - E[Y_0|X = 0] \\&= E[Y_1] - E[Y_0] \\&= ATE\end{aligned}$$

Experimental Inference VI

- ▶ This holds in experiments because of a physical process of randomization
- ▶ Units differ only in side of coin that was up
 - ▶ $X_i = 1$ only because $D_i = 1$
- ▶ Implications:
 - ▶ Covariate balance
 - ▶ Potential outcomes balanced and independent of treatment assignment
 - ▶ No confounding (selection bias)





Questions?

Experimental Analysis I

- ▶ The statistic of interest in an experiment is the *sample average treatment effect* (SATE)
- ▶ If our sample is representative, then this provides an estimate of the *population average treatment* (PATE)
 - ▶ Design-based random sampling
 - ▶ Model-based re-weighting
- ▶ This boils down to being a mean-difference between two groups:

$$SATE = \frac{1}{n_1} \sum Y_{1i} - \frac{1}{n_0} \sum Y_{0i}$$

Tidy Experimental Data

Sometimes it looks like this instead, which is bad:

unit	treatment	outcome0	outcome1
1	0	13	NA
2	0	6	NA
3	0	4	NA
4	0	5	NA
5	1	NA	3
6	1	NA	1
7	1	NA	10
8	1	NA	9

Tidy Experimental Data

An experimental data structure looks like:

unit	treatment	outcome
1	0	13
2	0	6
3	0	4
4	0	5
5	1	3
6	1	1
7	1	10
8	1	9

Computation of Effects I

- ▶ In practice we often estimate SATE using t-tests, ANOVA, or OLS regression
- ▶ These are all basically equivalent
- ▶ Reasons to choose one procedure over another:
 - ▶ Disciplinary norms
 - ▶ Ease of interpretation
 - ▶ Flexibility for >2 treatment conditions

Computation of Effects II

R:

```
t.test(outcome ~ treatment, data = data)  
lm(outcome ~ factor(treatment), data = data)
```

Questions?

Experimental Analysis II

- ▶ We don't just care about the size of the SATE. We also want to know whether it is significantly different from zero (i.e., different from no effect/difference)
- ▶ Thus we need to estimate the variance of the SATE
- ▶ The variance is influenced by:
 - ▶ Total sample size
 - ▶ Element variance of the outcome, Y
 - ▶ Relative size of each treatment group
 - ▶ (Some other factors)

Experimental Analysis III

Variance of the SATE

- ▶ $\widehat{Var}(SATE) = \frac{\widehat{Var}(\bar{Y}_0)}{n_0} + \frac{\widehat{Var}(\bar{Y}_1)}{n_1}$
 - ▶ Where $\widehat{Var}(Y_{0i})$ is control group variance
 - ▶ And $\widehat{Var}(Y_{1i})$ is treatment group variance

And the standard error of the variance is

- ▶ $\widehat{SE}_{SATE} = \sqrt{\frac{\widehat{Var}(\bar{Y}_0)}{n_0} + \frac{\widehat{Var}(\bar{Y}_1)}{n_1}}$

Intuition about Variance

- ▶ Bigger sample \rightarrow smaller SEs
- ▶ Smaller variance \rightarrow smaller SEs
- ▶ Efficient use of sample size:
 - ▶ When treatment group variances equal, equal sample sizes are most efficient
 - ▶ When variances differ, sample units are better allocated to the group with higher variance in Y

Statistical Power

Power analysis is used to determine sample size before conducting an experiment

	H_0 False ($ ATE > 0$)	H_0 True ($ATE = 0$)
Reject H_0	True positive	Type I Error
Accept H_0	Type II Error	True zero

True positive rate ($1 - k$) is power

False positive rate is the significance threshold (α)

Power calculation

- ▶ μ , Treatment group mean outcomes
- ▶ N , Sample size
- ▶ σ , Outcome variance
- ▶ α Statistical significance threshold
- ▶ Φ CDF of a sampling distribution.

$$\beta = \Phi \left(\frac{|\mu_t - \mu_c| \sqrt{N}}{2\sigma} - \Phi^{-1} \left(1 - \frac{\alpha}{2} \right) \right)$$

- ▶ β is Power, with distribution $[0,1]$.

Intuition about Power

Minimum detectable effect is the smallest effect we could detect given sample size, “true” ATE, variance of outcome measure, power $(1 - k)$, and α .

In essence: some non-zero effect sizes are not detectable by a study of a given sample size.

In underpowered study, we will be unlikely to detect true small effects. And most effects are small!²

²Gelman, A. and Weakliem, D. 2009. “Of Beauty, Sex and Power.” American Scientist 97(4): 310–16

Power in R

```
possible.ns <- seq(from = 100, to = 2000, by = 50) # The sample sizes we'll be considering
powers <- rep(NA, length(possible.ns)) # Empty object to collect simulation estimates
alpha <- 0.05 # Standard significance level
sims <- 500 # Number of simulations to conduct for each N

#### Outer loop to vary the number of subjects ####
for (j in 1:length(possible.ns)) {
  N <- possible.ns[j] # Pick the jth value for N


  significant.experiments <- rep(NA, sims) # Empty object to count significant experiments

  #### Inner loop to conduct experiments 'sims' times over for
  #### each N ####
  for (i in 1:sims) {
    Y0 <- rnorm(n = N, mean = 60, sd = 20) # control potential outcome
    tau <- 5 # Hypothesize treatment effect
    Y1 <- Y0 + tau # treatment potential outcome
    Z.sim <- rbinom(n = N, size = 1, prob = 0.5) # Do a random assignment
    Y.sim <- Y1 * Z.sim + Y0 * (1 - Z.sim) # Reveal outcomes according to assignment
    fit.sim <- lm(Y.sim ~ Z.sim) # Do analysis (Simple regression)
    p.value <- summary(fit.sim)$coefficients[2, 4] # Extract p-values (assuming
    # equal variance in treatment and control groups)
    significant.experiments[i] <- (p.value <= alpha) # Determine significance according to
    # p <= 0.05
  }

  powers[j] <- mean(significant.experiments) # store average success rate (power) for each N
}
```

Power in R


Power Calculation Different Sample Size ($\tau = 5$, $SD = 20$)




Factorial Designs

Factorial Designs

- ▶ The two-condition experiment is a stylized ideal
- ▶ An experiment can have any number of conditions
 - ▶ Up to the limits of sample size
 - ▶ More than 8–10 conditions is typically unwieldy
- ▶ Three “flavors”:
 - ▶ Multiple conditions in a single factor
 - ▶ Multiple fully crossed factors
 - ▶ Partially crossed (“fractional factorial”) designs
- ▶ Regression methods provide a generalizable tool for causal inference in such designs





Example

- ▶ How close do you feel to **your ethnic or racial group**?
- ▶ Some people have said that taxes need to be raised to take care of pressing national needs. How willing would you be to have your taxes raised to improve **education in public schools**?

Example

- ▶ How close do you feel to other Americans?
- ▶ Some people have said that taxes need to be raised to take care of pressing national needs. How willing would you be to have your taxes raised to improve education in public schools?

Example

- ▶ How close do you feel to **your ethnic or racial group**?
- ▶ Some people have said that taxes need to be raised to take care of pressing national needs. How willing would you be to have your taxes raised to improve **educational opportunities for minorities**?

Example

- ▶ How close do you feel to **other Americans**?
- ▶ Some people have said that taxes need to be raised to take care of pressing national needs. How willing would you be to have your taxes raised to improve **educational opportunities for minorities**?

2x2 Factorial Design

Condition	
Educ. for Minorities	Y_1
Schools	Y_0

2x2 Factorial Design

Condition	Americans	Own Race
Educ. for Minorities	$Y_{1,0}$	$Y_{1,1}$
Schools	$Y_{0,0}$	$Y_{0,1}$

Two ways to parameterize this

Dummy variable regression (i.e., treatment–control CATEs):

$$Y = \beta_0 + \beta_1 X_{0,1} + \beta_2 X_{1,0} + \beta_3 X_{1,1} + \epsilon$$

Interaction effects (i.e., treatment–treatment CATEs):

$$Y = \beta_0 + \beta_1 X_{1,1} + \beta_2 X_{2,1} + \beta_3 (X_{1,1} \times X_{2,1}) + \epsilon$$

Sample size for factorial designs

Condition	Conditions per factor	Total Conditions	n
1	2	2	400
1	3	3	600
1	4	4	800
2	2	4	800
2	3	6	1200
2	4	8	1600
3	3	9	1800
3	4	12	2400
4	4	14	3200

The examples assume power to detect a relatively small effect, but no consideration of multiple comparisons.

Considerations

- ▶ Factorial designs can quickly become unwieldy and expensive
- ▶ Need to consider what CATEs are of theoretical interest
 - ▶ Treatment-control, pairwise
 - ▶ Treatment-treatment, pairwise
 - ▶ Marginal effects, averaging across other factors
 - ▶ Comparison of merged conditions

Questions?